

CSC420 - Assignment 2

1. Implement Seam Carving

My implementation consists of two tables

- Cumulative Energy Table
 - Look up table used for the dynamic programming algorithm
 - Stores the cumulative energy of the shortest (least expensive) path from the top row to the bottom row
- Seam Path Table
 - Table that stores the position of the previous pixel in the shortest energy path along a seam

The Cumulative Energy at a given pixel is calculated as:

$$E(i, j) = G(i, j) + \min\{E(i - 1, j - 1), E(i - 1, j), E(i - 1, j + 1)\}$$

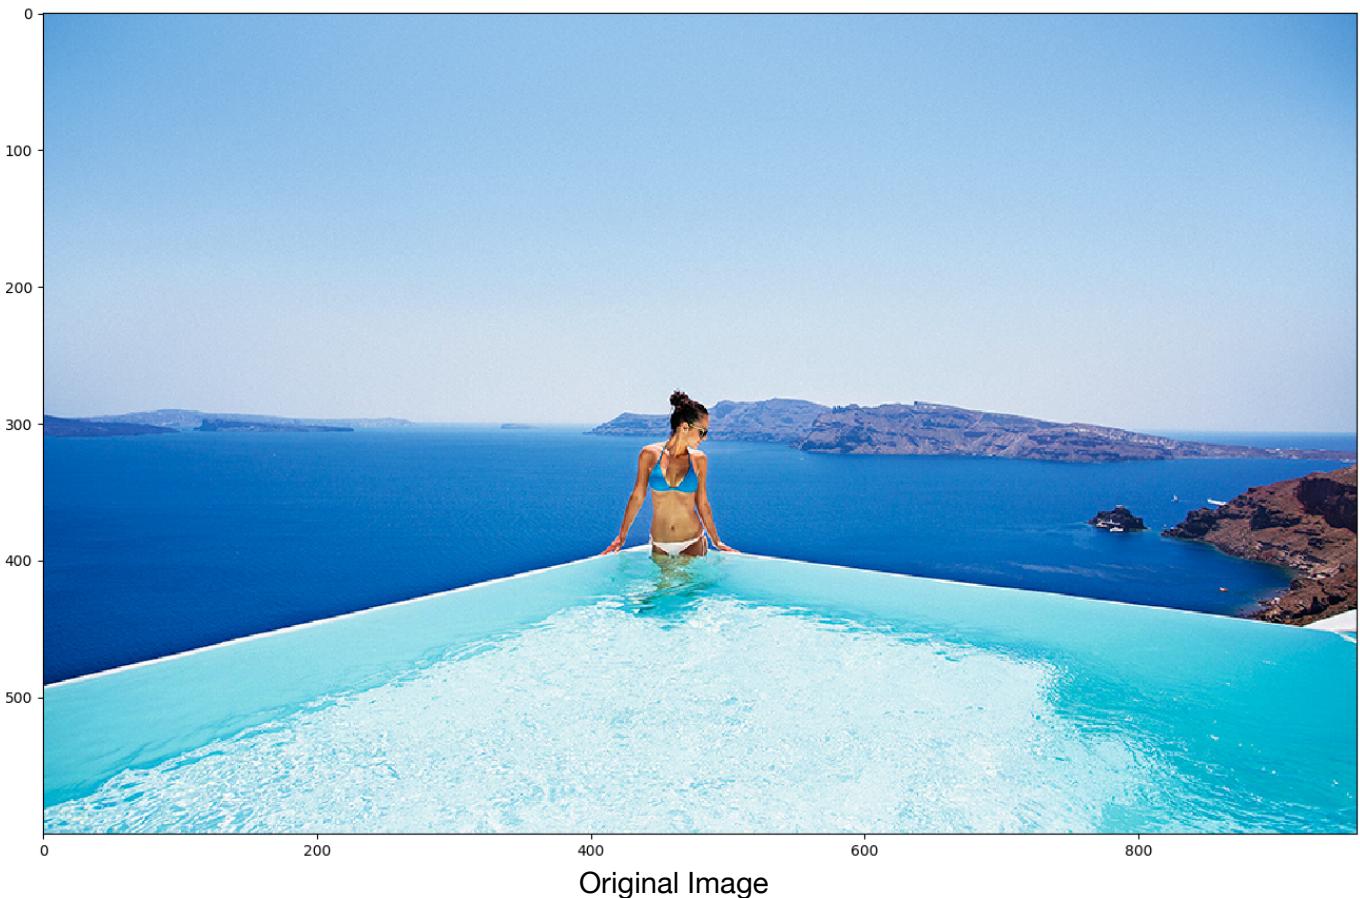
where E is the cumulative energy to point (i,j) and G is the gradient at that point

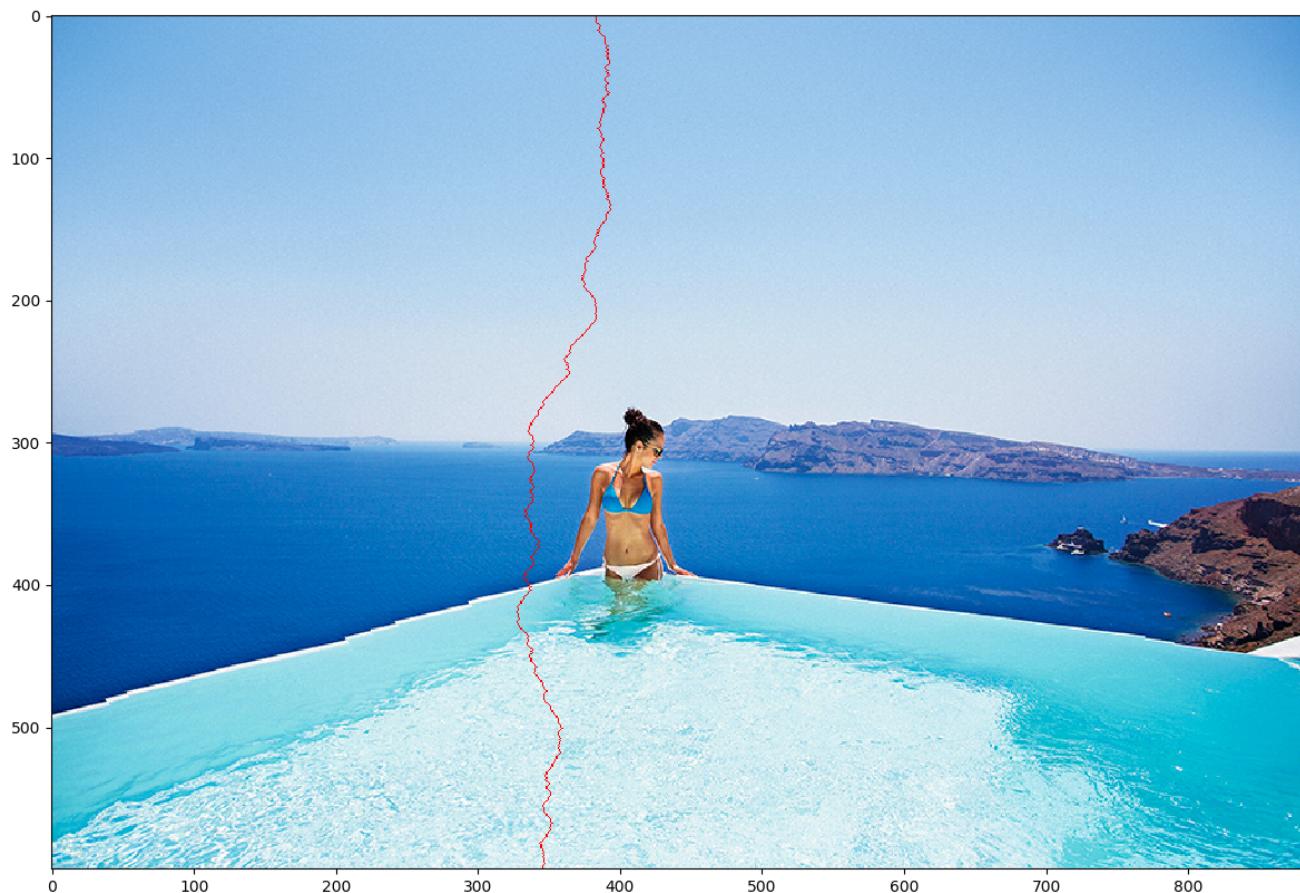
The Seam Path Table entry at a given pixel is the position of the previously traversed pixel with the lowest cumulative energy. It is calculated as:

$$P(i, j) = \arg \min_{(x, y)} \{E(x, y)\}, \text{ where } x = i - 1, \quad y \in \{j - 1, j, j + 1\}$$

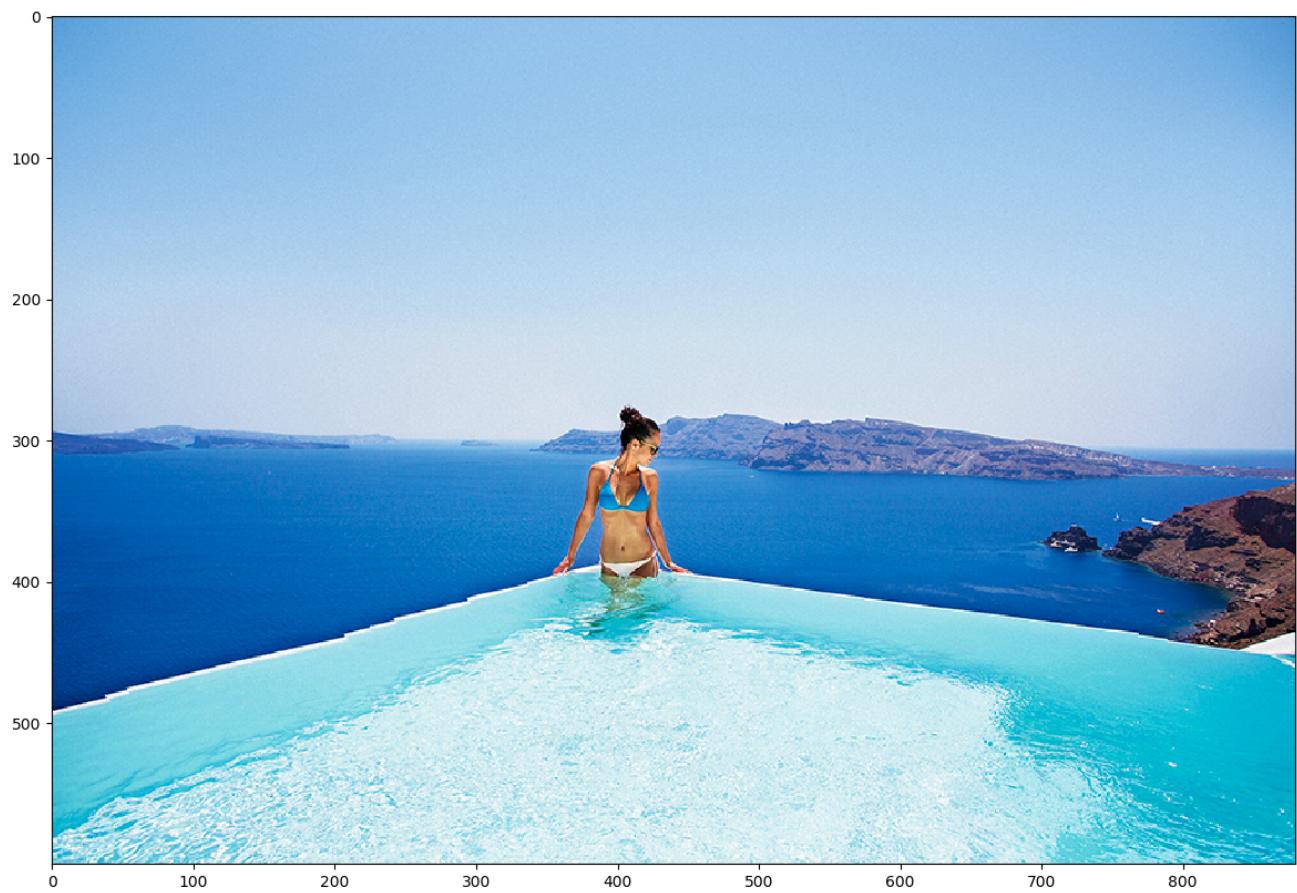
Seam Carving Example #1

The below image is seam carved 80 times to produce the final image.





80th Seam Removed



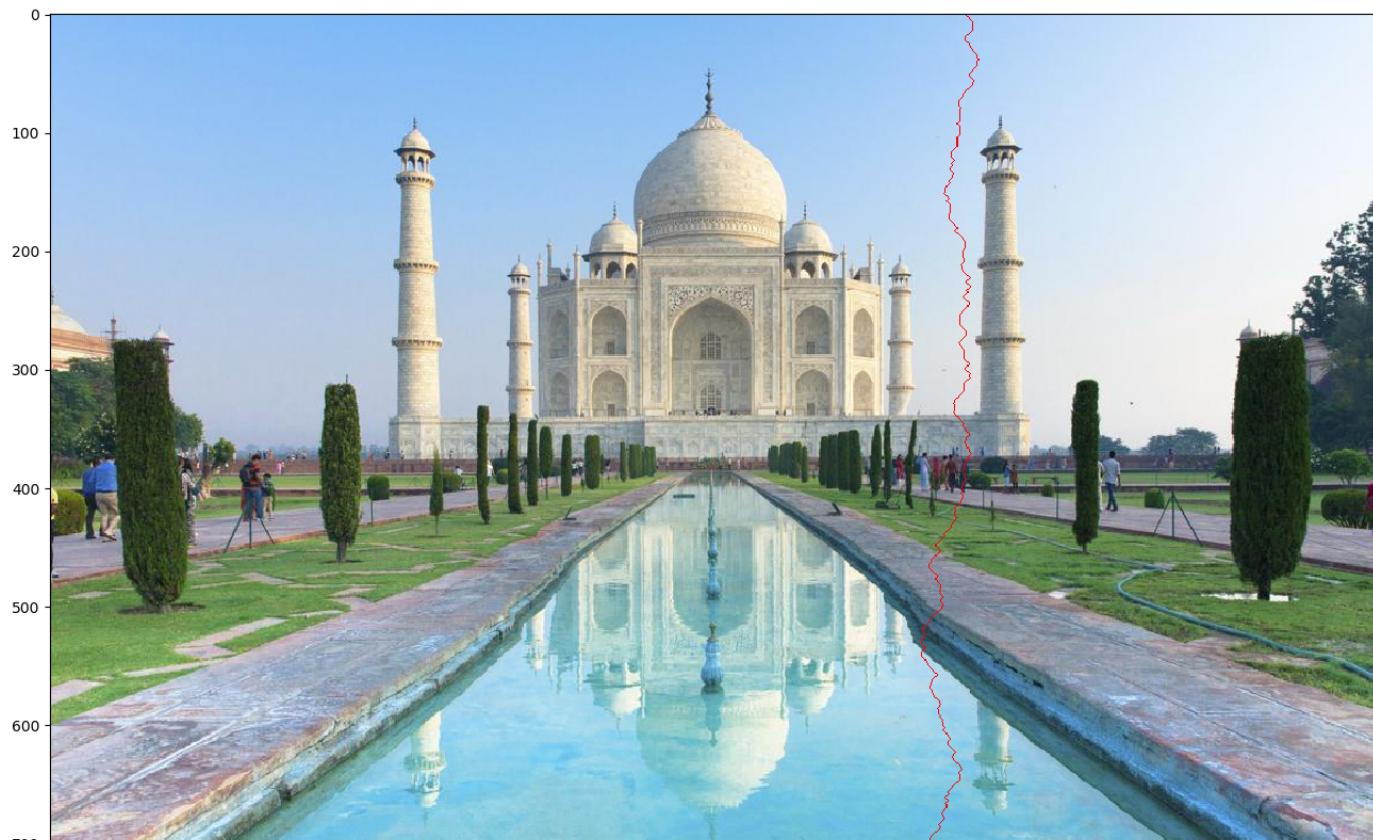
Final Seam Carved Image

Seam Carving Example #2

The below image is seam carved 80 times to produce the final image.



Original Image



80th Seam Removed



Final Seam Carved Image

2. Implement Harris Corner Detector

The implementation is straightforward and directly based off the lecture slides.

- My implementation uses the following parameters
 - gaussian window filter with standard deviation = 1, size = 5
 - alpha parameter = 0.04
 - threshold = 0.5 and 0.1

Key Algorithm Steps:

- 1) Create gaussian window filter (given a window size and standard deviation)
- 2) Find image gradient terms (I_{xx} , I_{xy} , I_{yy})
- 3) Compute elements of matrix M (M_{11} , M_{12} , M_{21} , M_{22})
- 4) Compute Cornerness ($R = (M_{11} * M_{22} - M_{12} * M_{21}) - \alpha * (M_{11} + M_{22})^2$) for each point
- 5) Threshold the Cornerness
- 6) Non-Maximum Supression

Harris Corner Example #1



Harris Corner Example #2



3. Implement SIFT Detector

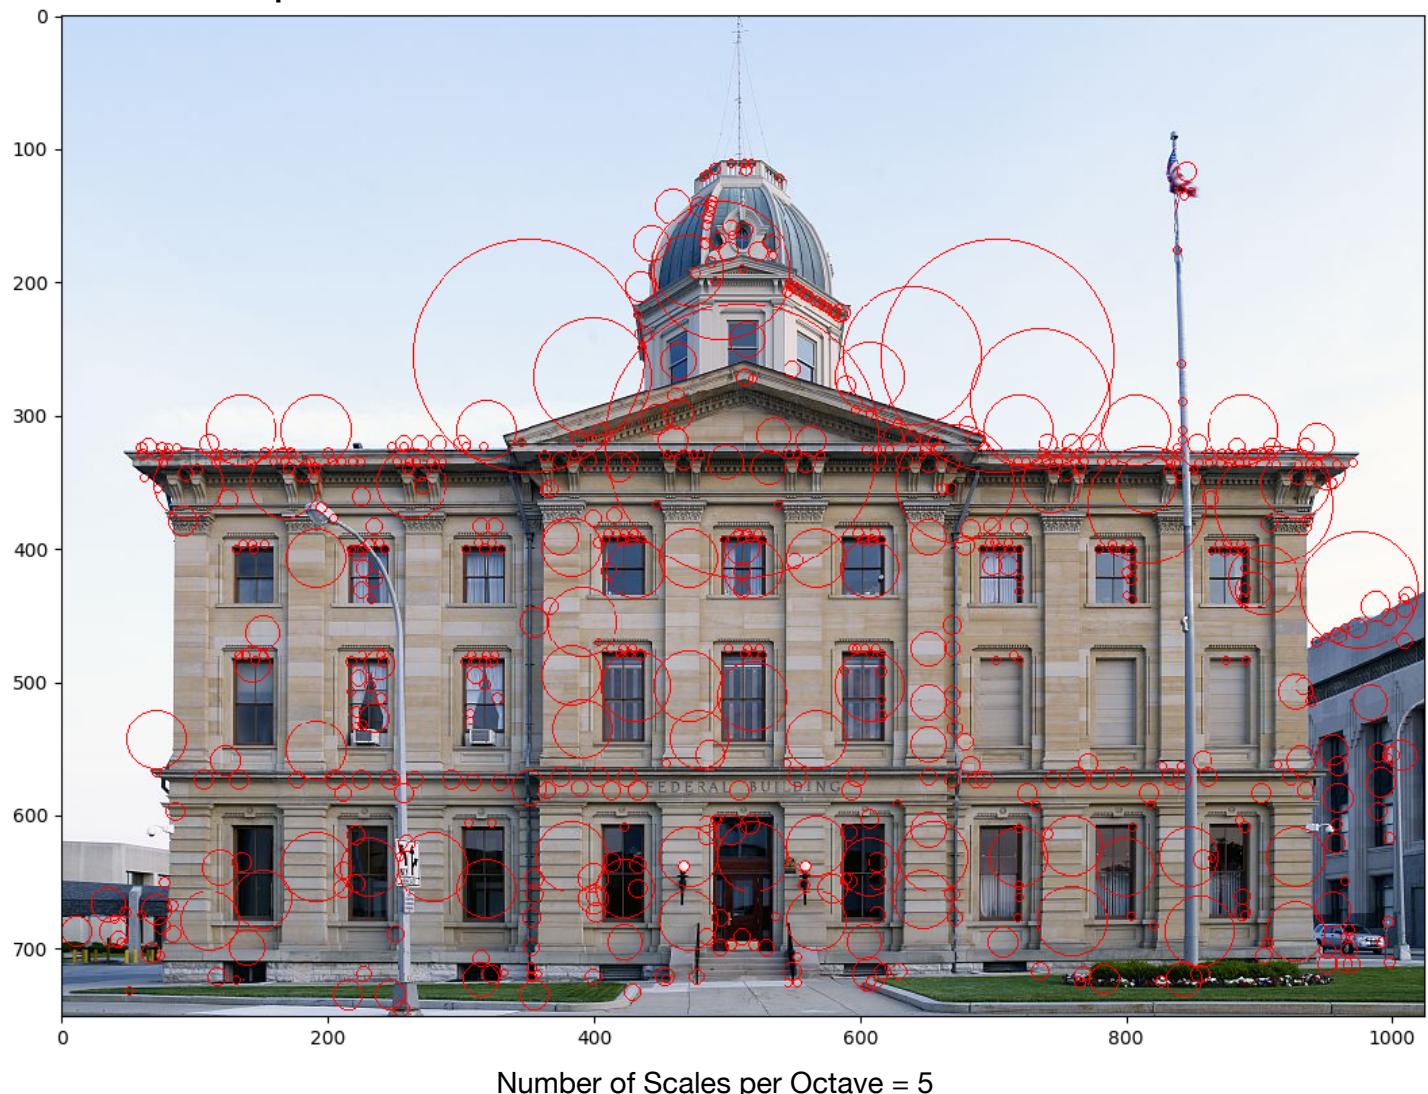
Most of the implementation is described in the lecture slides. However, there are two concepts that need explanation:

- Determining the location of a keypoint by upscaling the position of maxima at higher octaves
 - The pyramid_gaussian function takes every **odd** numbered index and discards the even numbered indices
 - We need to upscale by $2i + 1$ for each octave ($n > 1$)
 - A formula for cumulative upscaling at the nth octave is $2^{n-1}i + 2^{n-1} - 1$
 - this formula is off by 1 pixel for $n=1$. Since 1 pixel is inconsequential I disregarded this case when upscaling
 - Determining the radius of a blob/circle detected given the scale and octave of the maxima
 - Given a scale and octave we can calculate the equivalent standard deviation of the DoG/Laplacian.
 - The maxima of the Laplacian occurs when the zeros of the Laplacian are aligned with the circle
 - Therefore maximum response occurs at $r = \sigma\sqrt{2}$

The parameter values I used in my implementation are as follows:

- standard deviation used in the scale space = 1.6
- DoG threshold = 0.1

SIFT Detector Example #1



SIFT Detector Example #2



Appendix - Code

gradient.py (helper function)

```
import numpy as np
from scipy import ndimage

def gradient(image):
    sobel_x = np.array([
        [-1, 0, +1],
        [-2, 0, +2],
        [-1, 0, +1]
    ])
    sobel_y = np.array([
        [-1, -2, -1],
        [0, 0, 0],
        [+1, +2, +1]
    ])

    G_x = ndimage.correlate(image, sobel_x, mode='nearest')
    G_y = ndimage.correlate(image, sobel_y, mode='nearest')
    G = np.sqrt(G_x**2 + G_y**2)

    return G, G_x, G_y
```

seam_carving.py

```
import numpy as np
from skimage import io
from skimage.color import rgb2gray
from scipy import ndimage
from gradient import *

def minimum_seam(image):
    # FIND IMAGE GRADIENT
    G, G_x, G_y = gradient(image)
    height, width = G.shape

    # GENERATE DP LOOKUP TABLE
    E = np.full(G.shape, np.inf)      # Energy Table
    E[0] = G[0]
    P = np.empty_like(G, dtype=object) # Seam Path Table

    # POPULATE LOOK UP TABLES
    for i in range(1, height):
        for j in range(width):
            E_left = E[i-1,j-1] if j-1 >= 0 else np.inf
            E_right = E[i-1,j+1] if j+1 < width else np.inf
            E_center = E[i-1,j]

            # UPDATE ENERGY TABLE
            E[i,j] = G[i,j] + min(E_left, E_center, E_right)

    # UPDATE PATH TABLE
```

```

if E_left == min(E_left, E_center, E_right):
    P[i,j] = (i-1, j-1)
elif E_center == min(E_left, E_center, E_right):
    P[i,j] = (i-1, j)
elif E_right == min(E_left, E_center, E_right):
    P[i,j] = (i-1, j+1)

# FIND MINIMUM ENERGY PATH
min_i = height-1
min_j = np.argmin(E[height-1])
seam_root = (min_i, min_j)

# RETURN SEAM PATH AND SEAM ROOT INDEX
return P, seam_root

def seam_carving(image):
    image_gray = rgb2gray(image)
    P, seam_index = minimum_seam(image_gray)

    carve_image = np.zeros_like(image[:, :-1, :])
    seam_image = np.copy(image)

    while True:
        # Get index of the pixel on the seam
        i, j = seam_index[0], seam_index[1]

        # Draw Seam Line in Red
        seam_image[i,j] = (255,0,0)

        # Remove pixel
        carve_image[i,:,0] = np.delete(image[i,:,0], j)
        carve_image[i,:,1] = np.delete(image[i,:,1], j)
        carve_image[i,:,2] = np.delete(image[i,:,2], j)

        if i != 0: # Move to next pixel on seam path
            seam_index = P[i,j]
        else: # Exit loop after entire seam is traversed
            break

    return seam_image, carve_image

def main():
    image = io.imread('images/water.jpg', as_gray=False)

    carve_image = np.copy(image)
    for i in range(80):
        seam_image, carve_image = seam_carving(carve_image)
        print(carve_image.shape)

    io.imshow(image)
    io.show()

    io.imshow(seam_image)
    io.show()

    io.imshow(carve_image)
    io.show()

```

harris_corner_detector.py

```
import numpy as np
from skimage import io
from skimage import draw
from skimage.color import rgb2gray
from scipy import ndimage
from gradient import *

# RETURNS CORNERNESS MAP OF ONLY LOCAL MAXIMUM
def find_local_maximum(R):
    height, width = R.shape
    Corners = np.zeros_like(R)

    for i in range(height):
        for j in range(width):
            thresh = max(
                R[i-1,j-1] if i-1>=0 and j-1>=0 else 0,
                R[i-1,j]   if i-1>=0 else 0,
                R[i-1,j+1] if i-1>=0 and j+1<width else 0,
                R[i,j-1]   if j-1>=0 else 0,
                R[i,j+1]   if j+1<width else 0,
                R[i+1,j-1] if i+1<height and j-1>=0 else 0,
                R[i+1,j]   if i+1<height else 0,
                R[i+1,j+1] if i+1<height and j+1<width else 0
            )
            if R[i,j] >= thresh:
                Corners[i,j] = R[i,j]
    return Corners

def harris_corner_detector(image, window_size, stddev, thresh):
    image_gray = rgb2gray(image)

    # CREATE GAUSSIAN WINDOW FILTER
    if window_size%2==0: window_size=window_size+1
    k = int((window_size-1)/2)
    window = np.zeros((window_size, window_size))
    window[k, k] = 1
    window = ndimage.gaussian_filter(window, sigma=stddev)

    # FIND IMAGE GRADIENT TERMS
    G, I_x, I_y = gradient(image_gray)
    I_xx = I_x * I_x
    I_yy = I_y * I_y
    I_xy = I_x * I_y

    # COMPUTE M
    M_11 = ndimage.correlate(I_xx, window, mode='nearest')
    M_12 = M_21 = ndimage.correlate(I_xy, window, mode='nearest')
    M_22 = ndimage.correlate(I_yy, window, mode='nearest')

    # COMPUTE CORNERNESS R
    alpha = 0.04
    R = (M_11 * M_22 - M_12 * M_21) - alpha * (M_11 + M_22)**2

    # THRESHOLD R
    threshold = thresh
    R[R<threshold] = 0
```

```

# NON-LOCAL-MAXIMUM SUPPRESSION
Corners = find_local_maximum(R)

return R, Corners

def main():
    image = io.imread('images/building.jpg', as_gray=False)
    height, width, depth = image.shape

    R, Corners = harris_corner_detector(image, 5, stddev=1, thresh=0.5)

    # DRAW CORNERS ONTO IMAGE
    for i in range(height):
        for j in range(width):
            if Corners[i,j] > 0:
                rr, cc = draw.circle(i,j, 2, image.shape)
                image[rr,cc,:] = (255,0,0)

    io.imshow(image)
    io.show()

```

SIFT_keypoint_detector.py

```

import numpy as np
from skimage import io
from skimage import draw
from skimage.color import rgb2gray
from skimage.transform import pyramid_gaussian
from scipy import ndimage
from math import ceil

# n_scales: number of scales in the octave
# sigma: base standard deviation for each octave
def create_scale_space(octave, sigma, n_scales):
    shape = (octave.shape[0], octave.shape[1], n_scales)
    I = np.zeros(shape)
    k = 2**(1./(n_scales-1))

    for i in range(n_scales):
        stddev = (k**i) * sigma
        I[:, :, i] = ndimage.gaussian_filter(octave, sigma=stddev)

    return I

def difference_of_gaussians(I):
    shape = (I.shape[0], I.shape[1], I.shape[2]-1)
    D = np.zeros(shape)

    for i in range(1, I.shape[2]):
        D[:, :, i-1] = I[:, :, i] - I[:, :, i-1]

    return D

```

```

def check_local_extrema(D, i, j, k):
    target = D[i,j,k]
    for x in [i-1, i, i+1]:
        for y in [j-1, j, j+1]:
            for z in [k-1, k, k+1]:
                if D[x,y,z] > target:
                    return False
    return True

def find_extrema(D, thresh):
    n_scales = D.shape[2]
    height, width = D.shape[:2]
    extrema = []

    for k in range(1, n_scales-1):
        for i in range(1, height-1):
            for j in range(1, width-1):
                isLocalExtrema = check_local_extrema(D, i, j, k)
                # If a point is a local extrema add it to the list
                if isLocalExtrema:
                    # Threshold the point
                    if D[i,j,k] > thresh:
                        extrema.append((i,j,k))

    return extrema

def find_sift_keypoints(local_extrema, octave_num, sigma, n_scales):
    k = 2**(.1/(n_scales-1))
    keypoints = []

    for extrema in local_extrema:
        # Upsample the location and scale
        scaling_factor = (2 ** (octave_num-1))

        # i,j coordinate up-scaling
        location = (extrema[0] * scaling_factor + scaling_factor - 1,
                    extrema[1] * scaling_factor + scaling_factor - 1)

        # determining the scale at the keypoint
        scale = (k ** extrema[2]) * sigma * scaling_factor
        sift_keypoint = (location, scale)

        keypoints.append(sift_keypoint)

    return keypoints

# stddev: standard deviation used for gaussian smoothing of the
#         first scale in each octave
# n_scales: number of scales per octave
# thresh: threshold for the DoG extrema
def sift_keypoint_detector(image, stddev, n_scales, thresh):
    image_gray = rgb2gray(image)
    sift_keypoints = []
    oct_num = 1

    # COMPUTE GAUSSIAN PYRAMIDS
    pyramid = tuple(pyramid_gaussian(image_gray, multichannel=False))

    # iterate over each octave

```

for octave in pyramid:

```
# CREATE SCALE SPACE
I = create_scale_space(octave, sigma=stddev, n_scales=n_scales)

# COMPUTE DIFFERENCE OF GAUSSIANS
D = difference_of_gaussians(I)

# Modify D to make dealing with DoG easier
D = D**2      # Square to deal with only maxima
D_max = D.max()    # Scale to 1
D = D/D_max

# LOCATE LOCAL EXTREMA
local_extrema = find_extrema(D, thresh=thresh)

# STORE LOCAL EXTREMA POSITION AND SCALE IN RETURN LIST
keypoints = find_sift_keypoints(local_extrema, oct_num, stddev, n_scales)
sift_keypoints.extend(keypoints)

oct_num = oct_num + 1

return sift_keypoints
```

def main():

```
image = io.imread('images/building.jpg', as_gray=False)
```

```
keypoints = sift_keypoint_detector(image,
                                    stddev=1.6,
                                    n_scales=5,
                                    thresh=0.1)
```

Plot detected SIFT keypoints

for point in keypoints:

```
(i,j) = point[0]
stddev = point[1]
radius = ceil(2**(0.5) * stddev)
```

```
rr, cc = draw.circle_perimeter(i,j, radius, shape=image.shape)
image[rr,cc,:] = (255,0,0)
```

```
io.imshow(image)
io.show()
```